

MarkMatch marks the way you mark.

A marking assistant for college lecturers. The teacher marks six answers. The tool learns her marking from them, marks the rest of the class, quotes the line behind every mark, and nothing is final until she approves.

Open the live prototype

receipts for every mark ✓

MEASURED we ran it ourselves **SOURCE** we read it, link below **ASSUMED** not checked yet **UNKNOWN** we do not know

1 · PROBLEM

A lecturer cannot hand her marking to an AI, because it marks to its own standard, not hers.

Who has the problem

An assistant professor at a JNTUH-affiliated engineering college, marking the descriptive part of a mid-term for two sections.

ASSUMED A persona built from the exam rules. We have not interviewed a lecturer yet.

Pain 1: time

20–30 min

credited per undergraduate answer script by one state's order. Madras evaluators mark about 40 scripts a day.

SOURCE Govt of West Bengal order, 2019 · DT Next, June 2026

How often it comes round

2 × 20

Two mid-terms a semester per theory subject, each with a 20-mark descriptive part, plus assignments.

SOURCE JNTUH R22, read in one affiliated college's document

Pain 2: even trained humans do not agree

Two human graders marked the same 2,442 student answers. How often did they agree?

Gave the same mark		56.8%
Within one mark		78.0%

MEASURED by us on the Mohler dataset (a data structures course), 3 Oct 2026

So "the correct mark" depends on who is marking and how strict they are.

Why AI alone does not fix it

We pasted the same 8 real answers into ChatGPT in two separate chats. Class total out of 40:

ChatGPT, chat 1		21
ChatGPT, chat 2		15
Human grader 1		26
Human grader 2		24

MEASURED by us, 3 Oct 2026. One question, two chats: a small test.

The marks matched on 3 of 8 answers. One chat chose "partial credit", the other chose "strict". **Nobody told it to. It picked its own strictness.**

Why it happens

A marking scheme never covers every way a student can phrase an answer. Each marker, human or AI, fills the gaps with their own strictness.

Our reading of the two measurements above.

What a wrong mark costs

A student needs at least 14 of 40 internal marks to sit the semester-end exam. A teacher will not sign marks she cannot check.

SOURCE JNTUH R22, same college document

What we do not know

How many hours Indian college teachers really spend marking. Whether our own lecturers would use this. Real scripts are probably handwritten; we take typed answers today.

UNKNOWN We say so, and we ask two lecturers next.

AI marking already exists. In what we read, the trust step does not.

For each tool we state only what its own page says, or what we tested. We found no independent reviews, so we claim no weakness we did not read or reproduce.

OPTION	WHAT IT DOES	WHAT IS LEFT OPEN	BASIS
Marking by hand	The teacher reads every answer.	Slow, and two humans give the same mark only 56.8% of the time.	MEASURED
ChatGPT	Marks pasted answers and gives feedback.	The same answers got 21 and then 15. It marks one answer well; it is not repeatable.	MEASURED
Gemini in Google Classroom	Free. Its announcement describes rubric generation and suggested feedback on writing.	The announcement does not describe automatic marking.	SOURCE
CoGrader	AI first-pass grading with teacher review. Free for 100 submissions a month; handwriting on the \$15 plan.	We read nothing about where it falls short.	SOURCE
Eklavvy (India)	Scans handwritten exams, marks against a rubric. Clients include NMIMS and Sharda. Claims 31% less grading time.	Price not shown. Its page does not mention matching one teacher's standard.	SOURCE
University of Zurich platform	Highlights rubric matches and flags inconsistent grades after humans mark.	Deliberately shows no AI marks. Not offered publicly, as far as we could see.	SOURCE
Other hackathon builds (Evalyze)	Scan, compare to a model answer, show marks and a dashboard.	Its page reports no measured agreement with human graders.	SOURCE

We do not say nobody has built AI marking. Products exist and some are free. Parts of our idea also exist in research. What we did not find, in a limited search, is a tool a college lecturer can open that does the three things in the next section together.

Each loophole gets one answer.

THE AI PICKS ITS OWN STRICTNESS

Match my marking

She marks six answers. The tool marks the same six, counts how many match her, rewrites its notes and tries again. A new round is kept only if it matches her better.

SHE CANNOT SEE WHY A MARK WAS GIVEN

Every mark quotes the answer

Each mark shows **the exact line that earned it**. Code checks the line is really in the student's answer. If it is not, the mark is thrown out.

NOBODY SAYS HOW CLOSE THE TOOL IS

Proof on answers it has not seen

The tool marks ten answers whose real marks were held back, then shows how often it matched, beside how often two humans match each other.

THE TEACHER LOSES CONTROL

She approves, always

She approves the marking notes and every mark, and can change any of them. Nothing is final without her.

We do not claim better marks than ChatGPT. We claim her standard, checked evidence, and her signature.

A loop, four checks and a retry. We wrote every line.

HUMAN

Gives the question, her scheme, the typed answers, and her marks on six.

CODE

Checks the input. Splits "her six" from the unseen answers.

AGENT

Match my marking. The model marks the six; code counts matches with her; the model rewrites its notes; code keeps the round only if matches rise. At most two rounds.

HUMAN

Reads the notes, edits them, approves.

MODEL

Marks the rest, five answers per call. Every mark must carry a quote.

CODE

Hard checks: the quote is in the answer; the marks add up. A failure is asked again once, then sent to her unmarked.

CODE

Compares with her real marks on the unseen answers.

HUMAN

Edits any mark and approves the class.

Why we call it agentic, and how far

One step is an agent: it looks at where it disagrees with her, decides what to change, tries again, and is judged by code. Everything else is fixed code or a single model call, and the screen says which.

The model cannot overrule the checks. If the loop does not beat a single rewrite in our test, we will call it a checked pipeline and say so.

What we chose and why

- **Next.js, one app.** Screen and server together; one thing to deploy.
- **Our own engine, about 250 lines.** A framework would hide the steps a jury asks about.
- **Checks in code, not a second AI.** A text match is exact and free.
- **Gemini, with Groq as fallback.** Two free providers; if one refuses a call, the other takes it.
- **No database, no login.** Nothing is stored; marks stay in her browser tab.

A real run on the live site, with the real model.

✓ You mark six answers

HUMAN

✓ done · 6 of 6 entered

✓ Match my marking

AGENT

✓ Match my marking: matches you on 4 of 6

✓ You approve the scheme

HUMAN

✓ You approved the scheme

✓ Mark the unseen answers

MODEL

✓ Marked 10 answers

✓ Hard checks

CODE

✓ Hard checks: quote is in the answer · marks add up

✓ Compare with your real marks

CODE

✓ Compare with your real marks

8 You approve the marks

HUMAN

ROUND 0 • YOUR SCHEME AS TYPED

✓ KEPT

matches you on **4 of 6** total gap 2

ROUND 1 • NEW NOTES

✗ NOT KEPT · no better

matches you on **4 of 6** total gap 2

Give 5 points of c1 when answers state that nodes require additional memory for pointers.
Give 3 points of c2 when answers state that insertion and deletion are more complex or difficult.
Give 4 points by combining 1 point of c1 and 3 points of c2 when answers mention both complexity in pointer updates and increased memory overhead.

ROUND 2 • NEW NOTES

✗ NOT KEPT · no better

matches you on **3 of 6** total gap 3

Give 5 points of c1 when memory overhead of pointers is mentioned.
Give 3 points of c2 when insertion or deletion complexity is mentioned.
Give 1 additional point of c1 to answers that score 3 on c2 if they also mention managing

Nothing is final until you approve. 10 marks ready. Change any mark in its card first.

Approve marks

Proof

This tool vs you, same mark on 10 answers it had not seen

9 of 10 · scheme as typed within one mark: 9 of 10 · one question, this run

11.5 s · 8 model calls · 7,431 tokens · openai/gpt-oss-120b · fallback used · ₹0 (free tier)

How often two trained examiners agree, on 2,442 real answers:

same mark **56.8%**

within one mark **78.0%**

Mohler dataset · measured by us, 3 Oct

9 of 10

Same mark as the teacher on ten answers the tool had not seen, using her scheme as typed.

MEASURED one question, one run, 3 Oct 19:45

11.5 s

for the run: 8 model calls, 7,431 tokens, ₹0. The fallback provider took over when the first one ran out of free quota.

MEASURED shown on the screen above

4 of 6

Match my marking did not improve on her typed scheme in this run. The model wrote new notes twice; code found they matched her no better and **threw both rounds out**.

MEASURED The check worked. The tuning is what we are improving now.

FEATURE	STATE TONIGHT
Marks with a quoted line, checked by code	Working , tested on the live site
"Break it": plant a false quote	Working : thrown out on screen, fixed on the second try
Fallback to a second AI provider	Working : took over live once
Proof numbers on unseen answers	Working inside each run
Match my marking	Runs, but does not raise the match yet
Edit and approve marks	Built, still being tested
Recorded backup run · class mistake summary · handwritten scripts	Not built yet

How it is to use

One screen

No login, no dashboard. It opens on a real question with real student answers already loaded.

Her words

Mark, scheme, approve. Status is always a word with an icon, never colour alone.

Readable on a call

Light page, large type, text contrast above 7 to 1. Nothing depends on animation.

6 · SCALABILITY AND FEASIBILITY

A new subject costs six marks, not a retrained model.

What makes it scale

- **No training per subject.** A new question or a new teacher needs her scheme and six marks. This is how it is designed.
- **Nothing stored.** One stateless server route; no database to grow or guard.
- **Small calls.** Five answers per call, so a class of 60 is 12 calls per pass.
- **Cheap to run.** 16 answers took 8 calls and 7,431 tokens. **MEASURED**

What limits it today

- **Free quotas.** Gemini's free tier ran out after about two runs in a few minutes. Real use needs a paid key. **MEASURED**
- **Typed answers only.** Handwritten scripts are the next step, untested.
- **One subject tested.** All our numbers are from a data structures course in English.
- **Student data.** We used only a public dataset. Real scripts need the college's consent and a provider that does not train on them. **ASSUMED**

Who gains, and the number we will be judged on.

The lecturer

Reads the evidence for each mark and approves, instead of marking every answer from nothing.

Minutes saved per class: **UNKNOWN** until we time it with a real lecturer.

The student

The whole class is marked to one standard, and each mark comes with the line that earned it.

The department

Next: a summary of what most of the class got wrong, counted from the marks.

Not built yet.

First use case

Mid-term descriptive answers and assignments in one department: typed or transcribed short answers, one question at a time, with the lecturer approving. Our measure of success is agreement with her own marks on answers the tool has not seen, shown as "X of N" with the misses listed.

Short answers to hard questions.

▼ Why not just use ChatGPT?

It marks a single answer well. In our test the same eight answers got 21 and then 15, because it chose its own strictness each time. We fix whose standard is used, show checked evidence, and keep her approval.

▼ What stops the AI from inventing evidence?

Every mark needs a quote that is literally in the student's answer. Plain code tests that. The model cannot overrule it, and "Break it" shows it firing.

▼ What exactly is agentic here?

One step: Match my marking. It sees where it disagrees with the teacher, decides what to change, re-marks, and code keeps the change only if the match improves. The rest is fixed code or single model calls.

▼ Does Match my marking actually help?

Not yet in our live run: 4 of 6 before and after, so code discarded the new notes. On unseen answers the typed scheme scored 9 of 10. We report what we measured.

▼ Is the dataset in the model's training data?

Possibly; it is public and from 2011. That is why the judge can type an answer of their own, which the model has never seen.

▼ What about handwritten answer sheets?

Not supported today. It is the next thing we would test, with the teacher confirming the transcript before marking.

▼ Why your own engine and not LangChain or CrewAI?

It is about 250 lines: a loop, four checks and a retry. Every teammate can explain each step.

Sources

- Mohler ASAG dataset, CC-BY-4.0: huggingface.co/datasets/nkazi/MohlerASAG
- Govt of West Bengal order on evaluation hours, 22 Feb 2019: wbxpress.com
- DT Next, Madras University evaluation boycott, June 2026: dtnext.in
- JNTUH R22 internal evaluation, Scient Institute of Technology: scient.ac.in

- Gemini in Classroom announcement, 30 Jun 2025: [blog.google](https://blog.google/education/gemini-in-classroom/)
- CoGrader pricing: cograder.com/pricing
- Eklavvy AI answer sheet checking: eklavya.com
- University of Zurich AI-assisted grading platform: df.uzh.ch
- Evalyze, Devpost: devpost.com/software/evalyze

Student answers on the prototype are from the Mohler dataset (Mohler, Bunescu and Mihalcea, 2011), not from our college.